

Implementing Join

Join

1/87

DBMSs are engines to *store*, *combine* and *filter* information.

Filtering is achieved via selection and projection.

The *join* operation (\bowtie) is the primary means of *combining* information.

Because *join* is

- such an important operation in database applications/systems
- potentially very expensive to execute

many methods have been developed for its implementation.

(We use a running example to compare costs of the various join processing methods)

... Join

2/87

Types of join:

- simple equijoin (single-equality condition)

```
select * from R,S where R.i = S.j
```
- partial-match join (conjunction of equality conditions)

```
select * from R,S where R.a = S.b and R.c = S.d ...
```
- theta join (arbitrary expression as condition)

```
select * from R,S where R.a < S.b and R.c <> S.d ...
```

Focus on simple equijoin, since common in practice ($R.pk=S.fk$)

Join Example

3/87

Consider a university database with the schema:

```
create table Student(
  id      integer primary key,
  name    text, ...
);
create table Enrolled(
  stude   integer references Student(id),
  subj    text references Subject(code), ...
);
create table Subject(
  code    text primary key,
  title   text, ...
);
```

And the following request on this database:

List names of students in all subjects, arranged by subject.

... Join Example

4/87

The result of this request would look like:

Subj	Name
COMP1011	Chen Hwee Ling

```

COMP1011  John Smith
COMP1011  Ravi Shastri
...
COMP1021  David Jones
COMP1021  Stephen Mao
...
COMP3311  Dean Jones
COMP3311  Mark Taylor
COMP3311  Sashin Tendulkar

```

... Join Example

5/87

An SQL query to provide this information:

```

select E.subj, S.name
from   Student S, Enrolled E
where  S.id = E.stude
order  by E.subj, S.name;

```

And its relational algebra equivalent:

$$\text{Sort}[\text{subj}] (\text{Project}[\text{subj}, \text{name}] (\text{Join}[\text{id}=\text{stude}](\text{Student}, \text{Enrolled})))$$

The core of the query is the join $\text{Join}[\text{id}=\text{stude}](\text{Student}, \text{Enrolled})$

To simplify writing of formulae, $S = \text{Student}$, $E = \text{Enrolled}$.

... Join Example

6/87

Some database statistics:

Sym	Meaning	Value
r_S	# student records	20,000
r_E	# enrollment records	80,000
C_S	Student records/page	20
C_E	Enrolled records/page	40
b_S	# data pages in Student	1,000
b_E	# data pages in Enrolled	2,000

Also, in cost analyses below, N = number of memory buffers.

... Join Example

7/87

Out = $\text{Student} \bowtie \text{Enrolled}$ relation statistics:

Sym	Meaning	Value
r_{Out}	# tuples in result	80,000
C_{Out}	result records/page	80
b_{Out}	# data pages in result	1,000

Notes:

- r_{Out} ... one result tuple for each Enrolled tuple
- C_{Out} ... result tuples have only subj and name

Join via Cross-product

8/87

Join can be defined as a cross-product followed by selection:

$$\text{Join}[\text{Cond}](R,S) = \text{Select}[\text{Cond}](R \times S)$$

For the example query, could implement

$$\text{Join}[\text{id=stude}](\text{Student}, \text{Enrolled})$$

as

$$\text{Select}[\text{id=stude}](\text{Student} \times \text{Enrolled})$$

Cross-product contains $20,000 \times 80,000 = 1,600,000,000$ tuples.

... Join via Cross-product

9/87

For $\text{Temp} = (\text{Student} \times \text{Enrolled})$

I/O costs:

- size of Temp relation $r = 16 \times 10^8$ records
- assuming $C_{\text{Temp}}=16$, then $b_{\text{Temp}} = 10^8$
- Temp is written once, then scanned once
- total I/O = $10^8 \cdot (T_w + T_r)$

Assuming $T_w=T_r=0.01s$, this will take around 500 hours!

... Join via Cross-product

10/87

Because

- cross-products are infrequent in practice (except to describe join)
- cross-products are large (typically **much** larger than the final join result)

DBMSs do **not** implement join via cross-product.

DBMSs implement only join and provide cross-product as:

$$R \times S = \text{Join}[\text{true}](R,S)$$

or, in SQL

```
select * from R,S
```

Nested-Loop Join

Nested Loop Join

12/87

The simplest join algorithm:

- iteratively generates the cross-product
- checks join condition on each tuple

Algorithm to compute $\text{Join}[\text{Cond}](R,S)$:


```
for each tuple r in R {
  for each tuple s in S {
    if ((r,s) satisfies join condition) {
      add (r,s) to result
    }
  }
}
```

R is the *outer* relation; S is the *inner* relation.

... Nested Loop Join

13/87

Requires (at least) three memory buffers (2 input, 1 output).



... Nested Loop Join

14/87

Abstract algorithm for $Join[Cond](R,S)$ (with 3 memory buffers):

```

for each page of relation R {
  read into buffer rBuf
  for each page of relation S {
    read into buffer sBuf
    for each record r in rBuf {
      for each record s in sBuf {
        if ((r,s) satisfies Cond) {
          add combined(r,s) to OutBuf
          write Outbuf when full
        }
      }
    }
  }
}
  
```

... Nested Loop Join

15/87

Detailed algorithm for $Join[Cond](R,S)$ (with 3 memory buffers):

```

// rf: file for R, sf: file for S, of: output file
outp = 0; clearBuf(oBuf);
for (rp = 0; rp < nPages(rf); rp++) {
  readPage(rf, rp, rBuf);
  for (sp = 0; sp < nPages(sf); sp++) {
    readPage(sf, sp, sBuf);
    for (i = 0; i < nTuples(rBuf); i++) {
      rTup = getTuple(rBuf, i);
      for (j = 0; j < nTuples(sBuf); j++) {
        sTup = getTuple(sBuf, j);
        if (satisfies(rTup,sTup,Cond)) {
          rsTup = combine(rTup,sTup);
          addTuple(oBuf, rsTup);
          if (isFull(oBuf)) {
            writePage(of, outp++, oBuf);
            clearBuf(oBuf);
          }
        }
      }
    }
  }
}
  
```

... Nested Loop Join

16/87

The three-memory-buffer nested loop join requires:

- read all b_R pages of R once
- for each of page of R , read b_S pages of S

Cost = $b_R + b_R b_S$

If we use S as the outer relation in the join

Cost = $b_S + b_S b_R$

It is (slightly) better to use smaller relation as outer relation.

Nested Loop Join on Example

17/87

If Student is outer relation and Enrolled is inner:

$$\begin{aligned}\text{Cost} &= b_S + b_S b_E \\ &= 1,000 + 1,000 \times 2,000 = 2,001,000\end{aligned}$$

If Enrolled is outer relation and Student is inner:

$$\begin{aligned}\text{Cost} &= b_E + b_E b_S \\ &= 2,000 + 2,000 \times 1,000 = 2,002,000\end{aligned}$$

Cost of nested-loop join is too high (5 hours, if $T_r=0.01$ sec)

Implementing Join Better

18/87

Aims of effective join computation:

- generate only relevant tuples from the cross-product
- generate these tuples with minimal disk I/O

Range of costs for $Join(R,S)$


- worst case cost = $b_R + b_R b_S$ (nested loop join)
- best case cost = $b_R + b_S$ (read each page once)

Block Nested Loop Join

19/87

If at least b_R+2 memory buffers available:

- read the entire R relation into memory
- for each S page, check join condition on all (r, s) pairs



... Block Nested Loop Join

20/87

Algorithm for nested loop join with b_R+2 memory buffers:

```

read all of R's pages into memory buffers
for each page of relation S {
  read page into S's input buffer
  for each tuple s in S's buffer {
    for each tuple r in R's memory buffers {
      if ((r,s) satisfies JoinCond) {
        add (r,s) to output buffer
        write output buffer when full
      }
    }
  }
}
```

Note that R effectively becomes the inner relation in this scheme.

... Block Nested Loop Join

21/87

This method requires:

- read b_R pages of relation R into buffers
- while R is buffered, read b_S pages of S

$$\text{Cost} = b_R + b_S$$

Notes:

- minimal I/O cost, but considers all (r,s) pairs
- thus, requires $r_R.r_S$ checks of the join condition

... Block Nested Loop Join

22/87

Further performance improvements:

- must reduce number of R tuples matched against each S tuple
- use access method to find small set of R tuples matching S tuple

Example:

- each S joins with $k \ll r_R$ tuples of R
- R tuples are stored in sorted array of memory buffers
- for each S tuple, use binary search to find matching buffer
- scan around that buffer to find all matching (R,S) pairs
- requires approx $C_R.r_S$ checks of join condition

Block Nested Loop Join on Example

23/87

If ≥ 1002 memory buffers are available:

- read Student relation into memory
- scan Enrolled relation, computing join

$$\begin{aligned} \text{Cost} &= b_S + b_E \\ &= 1,000 + 2,000 = 3,000 \end{aligned}$$

This is considerably better than 10^6 (30 secs vs 5 hours).

But what if we have only N memory buffers, where $N < b_R$, $N < b_S$?

... Block Nested Loop Join on Example

24/87

In general case, read *outer* relation in runs of $N-2$ pages

```

for each run of N-2 pages from R {
  read N-2 of R's pages into memory buffers
  for each page of relation S {
    read page into S's input buffer
    for each tuple s in S's buffer do
      for each tuple r in R's memory buffers {
        if ((r,s) satisfies JoinCond)) {
          add (r,s) to output buffer
          write output buffer when full
        }
      }
    }
  }
}

```

... Block Nested Loop Join on Example

25/87

Block nested loop join requires

- read $\lceil b_R/N-2 \rceil$ runs from R
- for each run, scan b_S pages of S

$$\text{Cost} = b_R + b_S \cdot \lceil b_R / N - 2 \rceil$$

Notes:

- the final run will typically be "short" (i.e. $< N-2$ pages)
- unless index/hash is used, we still do $r_R.r_S$ tuple comparisons

... Block Nested Loop Join on Example

26/87

Costs for various buffer pool sizes:

N	Inner	Outer	#runs	Cost
22	Student	Enrolled	50	101,000
52	Student	Enrolled	20	41,000
102	Student	Enrolled	10	21,000
1002	Student	Enrolled	1	3,000
22	Enrolled	Student	100	102,000
52	Enrolled	Student	40	42,000
102	Enrolled	Student	20	22,000
1002	Enrolled	Student	2	4,000

Block Nested Loop Join in Practice

27/87

Why block nested loop join is very useful in practice ...

Many queries have the form

```
select * from R,S where r.i=s.j and r.x=k
```

This would typically be evaluated as

$$\text{Join } [i=j] ((\text{Sel}[r.x=k](R)), S)$$

If $|\text{Sel}[r.x=k](R)|$ is small \Rightarrow may fit in memory (in small #buffers)

Join Conditions and Methods

28/87

Nested loop join makes no assumptions about join conditions.

```
for each pair of tuples (r,s) {
    check join condition on (r,s)
    if satisfied, add to results
}
```

To improve join:

- reduce the number of tuple pairs considered
- but not easy to do for arbitrary join condition

As noted above, simple equijoin is a common join condition.

Thus, a range of other join algorithms has been developed specifically for equality join conditions.

Index Nested Loop Join

29/87

Most joins considered so far have a common problem:

- repeated scans of *entire* inner relation S are required

If there is an index on S, we can avoid such repeated scanning.

Consider $Join_{[R.i=S.j]}(R,S)$:

```

for each tuple r in relation R {
  use index to select tuples
    from S where s.j = r.i
  for each selected tuple s from S {
    add (r,s) to result
  }
}
```

(For ordered indexes (e.g. Btree), this also assists join conditions like $R.i < S.j$)

... Index Nested Loop Join

30/87

This method requires:

- one scan of R relation (b_R)
 - only one buffer needed, since we use R tuple-at-a-time
- for each *tuple* in R (r_R), one index lookup on S
 - cost depends on type of index and number of results
 - best case is when each $R.i$ matches few S tuples

Cost = $b_R + r_R \cdot Sel_S$ (Sel_S is the cost of performing a select on S).

... Index Nested Loop Join

31/87

For index lookup:

- cost of locating first matching tuple
 - for B+ trees, typically 2-4 page reads
 - for hashing, typically 1-2 page reads
- cost of finding other matching tuples
 - if clustered, typically 1-2 page reads
 - if unclustered, up to b_q page reads

Note: building an index "on the fly" to perform a join can be very cost-effective.

Index Nested Loop Join on Example

32/87

Case 1: $Join_{[id=stude]}(Student, Enrolled)$

- Student is outer and Enrolled is inner
- Enrolled has a clustered B+ tree index on stude field
- B+ tree has depth 3 (root + internal + leaf)
- most of the time, the four matching records are in a single page

$$\begin{aligned}
 \text{Cost} &= b_S + r_S \cdot btree_E \\
 &= 1,000 + 20,000 \times (3+1.01) = 80,000
 \end{aligned}$$

... Index Nested Loop Join on Example

33/87

Case 2: $Join_{[id=stude]}(Student, Enrolled)$

- Student is outer and Enrolled is inner
- Enrolled has an unclustered B+ tree index on stude field
- B+ tree has depth 3 (root + internal + leaf)
- assume worst case; matching records are all on different pages

$$\begin{aligned}
 \text{Cost} &= b_S + r_S \cdot btree_E \\
 &= 1,000 + 20,000 \times (3+4) = 150,000
 \end{aligned}$$

... Index Nested Loop Join on Example

34/87

Case 3: $Join_{[id=stude]}(Student, Enrolled)$

- Enrolled is outer and Student is inner
- Student is hashed on id field (e.g. linear hashing)
- there may be (short) overflow chains (e.g. 1.1 page reads/bucket)

$$\begin{aligned} \text{Cost} &= b_E + r_E \text{hash}_S \\ &= 2,000 + 80,000 \times 1.1 = 90,000 \end{aligned}$$

Optimised Index Nested Loop Join

35/87

Consider the following scenario for $Join_{[R.i=S.j]}(R, S)$:

- R.i is not a primary key (so many tuples have same R.i value)
- R is sorted on R.i (or could be efficiently sorted on R.i)
- each R.i value does not match very many tuples

Could save repeated index scans with the same R.i value

- cache results of index scan for $R.i=k$ in buffer
- if next R tuple also has $R.i=k$, re-use scan results

... Optimised Index Nested Loop Join

36/87

Abstract algorithm for optimised index nested loop join:

```
for each tuple r in relation R {
  if (prev == r.i)
    use selected tuples in buffer(s)
  else {
    use index to select tuples
    from S where s.j = r.i
    store selected tuples in buffer(s)
  }
  for each selected tuple s from S
    add (r,s) to result
  prev = r.i
}
```

Cost savings depend on repetition factor, #buffers, size of index scans

Sort-Merge Join

Sort-Merge Join

38/87

Basic approach:

- sort both relations on join attribute (reminder: $Join_{[R.i=S.j]}(R, S)$)
- scan together using merge to form result (r, s) tuples

Advantages:

- no need to deal with "entire" S relation for each r tuple
- deal with runs of matching R and S tuples

Disadvantages:

- cost of sorting both relations (relations may be sorted on join key?)


- some rescanning required when long runs of S tuples

... Sort-Merge Join

39/87

Method requires several cursors to scan sorted relations:

- r = current record in R relation
- s = start of current run in S relation
- ss = current record in current run in S relation



... Sort-Merge Join

40/87

Abstract algorithm for merge phase of $Join_{[R.i=S.j]}(R,S)$:

```

r = first tuple in R
s = first tuple in S
while (r != eof and s != eof) {
    // align cursors to start of next common run
    while (r != eof and r.i < s.j) { r = next tuple in R }
    while (s != eof and r.i > s.j) { s = next tuple in S }
    // scan common run, generating result tuples
    while (r != eof and r.i == s.j) {
        ss = s // set to start of run
        while (ss != eof and ss.j == r.i) {
            add (r,s) to result
            ss = next tuple in S
        }
        r = next tuple in R
    }
    s = ss // start search for next run
}

```

Sidetrack: Iterators

41/87

Sort-merge join implementation is simplified by use of iterators.

- iterators give the appearance of tuple-at-a-time
- even when the underlying data is page-by-page
- and even in the presence of auxiliary index structures

Typical usage of iterator:

```

Iterator iter; Tuple tup;
iter = startScan("Rel", "i=5");
while ((tup = nextTuple(iter)) != NULL) {
    process(tuple);
}
endScan(iter);

```

... Sidetrack: Iterators

42/87

```

typedef struct {
    File    inf; // input file
    Buffer buf;  // buffer holding current page
    int     curp; // current page during scan
}

```

```

    int    curr; // index of current record in page
} Iterator;

// simple linear scan; no condition
Iterator *startScan(char *relName) {
    Iterator *iter = malloc(sizeof(Iterator));
    iter->inf = openFile(fileName(relName),READ);
    iter->curp = 0;
    iter->curr = -1;
    readPage(iter->inf, iter->curp, iter->buf);
}

```

... Sidetrack: Iterators

43/87

```

Tuple nextTuple(Iterator *iter) {
    // check if reached end of current page
    if (iter->curr == nTuples(iter->buf)-1) {
        // check if reached end of data file
        if (iter->curp == nPages(iter->inf)-1)
            return NULL;
        iter->curp++;
        iter->buf = readPage(iter->inf, iter->curp);
        iter->curr = -1;
    }
    iter->curr++;
    return getTuple(iter->buf, iter->curr);
}
// curp and curr hold indexes of most recently read page/record

```

... Sidetrack: Iterators

44/87

```

TupleID scanCurrent(Iterator *iter) {
    // form TupleID for current record
    return iter->curp + iter->curr;
}

void setScan(Iterator *iter, int page, int rec) {
    assert(page >= 0 && page < nPages(iter->inf));
    if (iter->curp != page) {
        iter->curp = page;
        readPage(iter->inf, iter->curp, iter->buf);
    }
    assert(rec >= 0 && rec < nTuples(iter->buf));
    iter->curr = rec;
}

void endScan(Iterator *iter) {
    closeFile(iter->buf);
    free(iter);
}

```

Sort-Merge Join

45/87

Concrete algorithm using iterators:

```

Iterator *ri, *si; Tuple rup, stup;

ri = startScan("SortedR");
si = startScan("SortedS");
while ((rtup = nextTuple(ri)) != NULL
    && (stup = nextTuple(si)) != NULL) {
    // align cursors to start of next common run
    while (rtup != NULL && rtup.i < stup.j)
        rtup = nextTuple(ri);
    if (rtup == NULL) break;
    while (stup != NULL && rtup.i > stup.j)

```

```

        stup = nextTuple(si);
    if (stup == NULL) break;
        // must have (r.i == s.j) here
...

```

... Sort-Merge Join

46/87

```

...
// remember start of current run in S
TupleID startRun = scanCurrent(si);
// scan common run, generating result tuples
while (rtup != NULL && rtup.i == stup.j) {
    while (stup != NULL and stup.j == rtup.i) {
        addTuple(outbuf, combine(rtup, stup));
        if (isFull(outbuf)) {
            writePage(outf, outp++, outbuf);
            clearBuf(outbuf);
        }
        stup = nextTuple(si);
    }
    rtup = nextTuple(ri);
    setScan(si, startRun);
}
}

```

... Sort-Merge Join

47/87

Buffer requirements:

- for sort phase:
 - as many as possible (remembering that cost is $O(\log_{\#Bufs})$)
 - if insufficient buffers, sorting cost can dominate
- for merge phase:
 - one output buffer for result
 - one input buffer for relation R
 - (preferably) enough buffers for longest run in S

... Sort-Merge Join

48/87

Cost of sort-merge join.

Step 1: sort each relation (if not already sorted):

- Cost = $2 \cdot b_R (1 + \log_{N-1}(b_R/N)) + 2 \cdot b_S (1 + \log_{N-1}(b_S/N))$
(where N = number of memory buffers)

Step 2: merge sorted relations:

- if every run of values in S fits completely in buffers, merge requires single scan, Cost = $b_R + b_S$
- if some runs in of values in S are larger than buffers, need to re-scan run for each corresponding value from R

Sort-Merge Join on Example

49/87

Case 1: $Join[id=stude](Student, Enrolled)$

- *Student* and *Enrolled* already sorted on *id#*
- memory buffers $N=4$; all runs are of length < 2

Cost = $b_S + b_E = 3,000$ (i.e. minimal cost)

... Sort-Merge Join on Example

50/87

Case 2: $\text{Join}_{[id=stude]}(\text{Student}, \text{Enrolled})$

- relations are not sorted on $id\#$
- memory buffers $N=32$; all runs are of length < 30

$$\begin{aligned} \text{Cost} &= \text{sort}(S) + \text{sort}(E) + b_S + b_E \\ &= b_S \lceil \log_{30} b_S \rceil + b_E \lceil \log_{30} b_E \rceil + b_S + b_E \\ &= 1,000 \times 3 + 2,000 \times 3 + 1,000 + 2,000 \\ &= 12,000 \end{aligned}$$

... Sort-Merge Join on Example

51/87

Case 3: $\text{Join}_{[id=stude]}(\text{Student}, \text{Enrolled})$

- *Student* and *Enrolled* already sorted on $id\#$
- memory buffers $N=3$ (*S* input, *E* input, output)
- one-quarter of the "runs" in *E* span two pages
- there are no "runs" in *S*, since $id\#$ is a primary key

Cost depends on which relation is outer and which is inner.

... Sort-Merge Join on Example

52/87

Case 3 (continued) ...

If *E* is outer relation:

- $\text{Cost} = b_E + b_S = 3,000$

If *S* is outer relation:

- one-quarter of *E* runs require two page reads
- each *E* run is processed once for matching *S.id* value
- $\text{Cost} = b_S + b_E + r_S/4 = 8,000$

Sidetrack 2: More on Iterators

53/87

Above description of iterators:

- involved simple scan of a single table
- with no condition to select tuples

In the general case, an iterator involves:

- one (selection) or two (join) tables
- with a condition to determine relevant tuples

A typical SQL query involves many iterators

- one for each relational operator in query plan
- connected in a demand-driven network of query nodes

... Sidetrack 2: More on Iterators

54/87

Requires a more general definition of execution state:

```
typedef struct {
    Oper    op;      // operation (sel, sort, join, ...)
    Reln    r1;      // first relation
    Reln    r2;      // second relation (if any)
    Buffer   *bufs;   // buffers used by operation
```

```

int    curp1; // index of current page for r1
int    curr1; // index of current record in page
int    curp2; // index of current page for r2
int    curr2; // index of current record in page
Cond   cond;  // condition for choosing tuple(s)
} Iterator;

```

For PostgreSQL details, see [include/nodes/execnodes.h](#)

Hash Join

Hash Join

56/87

Basic idea:

- use hashing as a technique to partition relations
- to avoid having to consider all pairs of tuples

Requires sufficient memory buffers

- to hold substantial portions of partitions
- (preferably) to hold largest partition of outer relation

Other issues:

- works only for equijoin $R.i=S.j$ (but this is a common case)
- susceptible to data skew (or poor hash function)

Variations: *simple*, *grace*, *hybrid*.

Simple Hash Join

57/87

Basic approach:

- hash the inner relation R into memory buffers (build)
- scan the outer relation S , using hash to search (probe)
 - if $R.i=S.j$, then $h(R.i)=h(S.j)$ (hash to same buffer)
 - only need to check one memory buffer for each S tuple


Makes the assumption: whole of S hashes into memory

- requires R to be smaller than memory buffers
- requires a uniform hash function (no overflows)

... Simple Hash Join

58/87

Data flow:



... Simple Hash Join

59/87

Algorithm for ideal simple hash join $Join[R.i=S.j](R,S)$:

```

for each tuple r in relation R
  { insert r into buffer[h(R.i)] }
for each tuple s in relation S {
  for each tuple r in buffer[h(S.j)] {
    if ((r,s) satisfies join condition) {
      add (r,s) to result
    }
  }
}

```

Cost = $b_R + b_S$ (minimum possible cost)

... Simple Hash Join

60/87

Consider that we have N buffers available.

If $b_R \leq N-2$ buffers, no need to hash (use nested loop).

In practice, size of hash table $b_{hR} > b_R$ (e.g. data skew)
 \Rightarrow hash table for R is even less likely to fit in memory

Can be handled by a variation on above algorithm:

- scan R , making hash table with $N-2$ buffers
- once hash table built, scan S (standard probe phase)
- if more R tuples, build new table and repeat

... Simple Hash Join

61/87

Algorithm for realistic simple hash join $Join[R.i=S.j](R,S)$:

```

for each tuple r in relation R {
  if (buffer[h(R.i)] is full) {
    for each tuple s in relation S {
      for each tuple rr in buffer[h(S.j)] {
        if ((rr,s) satisfies join condition) {
          add (rr,s) to result
        }
      }
    }
    clear all hash table buffers
  }
  insert r into buffer[h(R.i)]
}

```

Note: requires multiple passes over the S relation.

... Simple Hash Join

62/87

Cost depends on N and on properties of data/hash.

Worst case:

- $h(i)=k$ so read only C_R tuples before hash table "full"
- each hash table for R occupies one buffer with C_R tuples
- degenerates to nested-loop-with-3-buffers case $\Rightarrow b_R + b_R b_S$

Best case:

- perfect uniform distribution of hash values
- each hash table of R holds $(N-2)C_R$ tuples from $N-2$ pages
- number of hash tables built = $n_{hR} = \lceil b_R / (N-2) \rceil$
- read all of S for each hash table $\Rightarrow b_R + n_{hR} \cdot b_S$

Grace Hash Join

Basic approach:

- partition both relations on join attribute using hashing
- scan through corresponding pairs of partitions to form results

Similar approach to sort-merge join, except:


- sort-merge: partitioning achieved by sorting (runs)
- hash: partitioning achieved by hashing

Requires enough buffer space to hold largest partition of inner relation.

... Grace Hash Join

64/87

Partition phase:




This is applied to each relation R and S .

... Grace Hash Join

65/87

Probe/join phase:



The second hash function ($h2$) simply speeds up the matching process. Without it, would need to scan entire R partition for each record in S partition.

... Grace Hash Join

66/87

Abstract algorithm for $Join_{R.i=S.j}(R,S)$:

```
// assume h(val) generates [0..N-2]
// assume h2(val) generates [0..N-3]

// Partition phase (each relation -> N-1 partitions)
// 1 input buffer, N-1 output buffers

for each tuple r in relation R
    add r to partition h(r.i) in output file R'
```



```

for each tuple s in relation S
    add s to partition h(s.j) in output file S'
...

```

... Grace Hash Join

67/87

Abstract algorithm for $Join[R.i=S.j](R,S)$ (cont.)

```

// Probe/join phase
// 1 input buffer for S, 1 output buffer
// N-2 buffers to build hash table for R partition

for each partition p = 0 .. N-2 {
    // Build in-memory hash table for partition p of R'
    for each tuple r in partition p of R'
        insert r into buffer h2(r.i)

    // Scan partition p of S', probing for matching tuples
    for each tuple s in partition p of S' {
        b = h2(s.j)
        for all matching tuples r in buffer b
            add (r,s) to result
    }
}

```

... Grace Hash Join

68/87

Concrete algorithm for partitioning:

```

Buffer iBuf, oBuf[N-1];
File inf, outf[N-1]; char rel[100];
int i, r, h, ip, op[N-1]; Tuple tup;
for (i = 0; i < N-1; i++) {
    clearBuf(oBuf[i]); op[i] = 0;
    rel = sprintf("%s%d", "Rel", i);
    outf[i] = openFile(fileName(rel), WRITE);
}
inf = openFile(fileName("Rel"), READ);
for (ip = 0; ip < nPages(inf); ip++) {
    iBuf = readPage(inf, ip);
    for (r = 0; r < nTuples(iBuf); r++) {
        tup = getTuple(iBuf, r);
        h = hash(tup.i, N-1);
        addTuple(oBuf[h], tup);
        if (isFull(oBuf[h])) {
            writePage(outf[h], op[h]++, oBuf[h]);
            clearBuf(oBuf[h]);
        }
    }
}

```

... Grace Hash Join

69/87

Cost of grace hash join:

- #pages in all partition files of $Rel \approx b_{Rel}$ (maybe slightly more)
- partition relation R ... Cost = $b_R.T_r + b_R.T_w = 2b_R$
- partition relation S ... Cost = $b_S.T_r + b_S.T_w = 2b_S$
- probe/join requires one scan of each (partitioned) relation
Cost = $b_R + b_S$
- all hashing and comparison occurs in memory $\Rightarrow \approx 0$ cost

Total Cost = $3(b_R + b_S)$

... Grace Hash Join

70/87

The above cost analysis assumes:

- every partition of R fits in memory buffers at once

We achieve this situation if:

- data has uniform distribution
- hash function gives uniform distribution
(all partitions are similar size)
- we have $N-1 \geq \lceil \sqrt{b} \rceil$ memory buffers
(giving $N-1$ partitions, each with $\approx b_R/(N-1)$ pages)

... Grace Hash Join

71/87

Possibilities for dealing with "over-long" partitions of R

- handle each over-long partition via scanning
 - requires over-long partitions to be scanned multiple times
 - essentially, such partitions are treated via nested loop join
- apply hash join recursively to over-long partitions
 - increases i/o by needing to partition parts of file multiple times
- use a different hash function with better distribution properties
 - but difficult to find such hash functions "on the fly"
- use the relation with the best partitioning as the "outer" relation

Grace Hash Join on Example

72/87

For the example $Join[id=stude](Student, Enrolled)$:

- assume that we have a good hash function and $N = \sqrt{1000} = 32$

$$\begin{aligned} \text{Cost} &= 3(b_S + b_E) \\ &= 3(1,000 + 2,000) = 9,000 \end{aligned}$$

Hybrid Hash Join

73/87

An optimisation if we have $\sqrt{b_R} < N < b_R+2$

- create k partitions using N buffers where $k \ll N$
- with grace join, would use k output buffers (one per partition)
- what to do with $N-k$ remaining buffers? (ignore input buffer)
- use them to hold m partitions of R in memory (no disk writes)
- other partitions are handled as before (using $k-m$ output buffers)

When we come to scan and partition S relation

- any tuple with hash in range $0..m-1$ can be resolved
- other tuples are written to one of $k-m$ partition files for S

Final phase is same as grace join, but with only $k-m$ partitions.

... Hybrid Hash Join

74/87

Some observations:

- for k partitions, each partition has expected size $\text{ceil}(b_R/k)$
- holding m partitions in memory needs $m \times \text{ceil}(b_R/k)$ buffers
- since we have $k-m$ output buffers, we must have $m b_R/k + (k-m) \leq N$
- for every partition/block held in memory, we save on disk i/o
- saving is $m/k \times 2(b_R + b_S)$

Other notes:

- if $N = b_R+2$, using block nested loop join is simpler
- cost depends on N (but less than grace hash join)

... Hybrid Hash Join

75/87

Need to choose appropriate m and k to minimise cost

- base cost: $3 \times (b_R + b_S)$ (grace join)
- i/o saving: $m/k \times 2(b_R + b_S)$
- constraint: $mb_R/k + (k-m) \leq N$


Approach to maximise saving:

- have one large in-memory partition ($m = 1$)
- use as many as possible of N buffers for partition
- use as few output buffers as possible (minimise k)

... Hybrid Hash Join

76/87


Data flow for hybrid hash join (partitioning R):



... Hybrid Hash Join

77/87

Data flow for hybrid hash join (partitioning S):



After this, proceed as for grace hash join.

... Hybrid Hash Join

78/87

Cost of hybrid hash join:

- assume: large N total buffers, m partitions in memory, k partitions on disk
- read both tables: $b_R + b_S$
- total partitions for each table: $m+k$
- assuming uniform hashing, #pages in each R partition $P_R = \text{ceil}(b_R/(m+k))$
- assuming uniform hashing, #pages in each S partition $P_S = \text{ceil}(b_S/(m+k))$
- in Pass 1, $k \cdot P_R + k \cdot P_S$ pages written to disk partitions
- all joining of m in-memory partitions is handled in memory
- in Pass2, $k \cdot P_R + k \cdot P_S$ pages read back from disk partitions

$$\begin{aligned} \text{Cost} &= b_R + b_S + k \cdot P_R + k \cdot P_S + k \cdot P_R + k \cdot P_S \\ &= b_R + b_S + 2 \cdot k \cdot (P_R + P_S) \end{aligned}$$

$$= b_R + b_S + 2 * k * (\text{ceil}(b_R/(m+k)) + \text{ceil}(b_S/(m+k)))$$

How to determine k :

- set $m=1$ and so size of partition $\cong N \Rightarrow k \cong b_R/N$
- need to ensure that $\lceil b_R/k \rceil + k \leq N$ (allowing for input buffer)
- choose k close to b_R/N but satisfying constraint

Hybrid Hash Join on Example

79/87

Case 1: $N = 100$ buffers, $b_R = 1000$

- $k = 10 \Rightarrow 1000/10 + 10 = 110$ buffers; not less than 100
- $k = 12 \Rightarrow 1000/12 + 12 = 96$ buffers
- Cost = $(3-2/12) \cdot (1000+2000) = 8500$

Case 2: $N = 200$ buffers, $b_R = 1000$

- $k = 5 \Rightarrow 1000/5 + 5 = 205$ buffers; not less than 200
- $k = 6 \Rightarrow 1000/6 + 6 = 173$ buffers
- Cost = $(3-2/6) \cdot (1000+2000) = 8000$

Case 3: $N = 502$ buffers, $b_R = 1000$

- $k = 2 \Rightarrow 1000/2 + 2 = 502$ buffers
- Cost = $(3-2/2) \cdot (1000+2000) = 6000$

Pointer-based Join

80/87

Conventional join algorithms set up $R \leftrightarrow S$ connections via attribute values.

Join could be performed faster if direct connections already existed.

- in OODBMSs, they generally already exist in the form of object references (oids)
- in RDBMSs, they could be introduced via extra rid attributes

Such a modification to conventional RDBMS structure would be worthwhile:

- if we know in advance what kind of joins will be required
- adding the extra rid attributes into tuples is feasible

... Pointer-based Join

81/87

The basic idea for pointer-based join is:

```
for each tuple r in relation R {
  for each rid associated with r {
    fetch tuple s from S via rid
    add (r,s) to result relation
  }
}
```

Often, each R tuple is associated with only one rid, so the inner loop is not needed.

... Pointer-based Join

82/87

The advantage over value-based joins:

- rather than find S tuples via value-based lookup (e.g. hashing, index)
- we find S tuples by direct fetch with rid (much faster per tuple)
- requires no assumption about sorted-ness of relations
- does not require large numbers of buffers

The (potential) disadvantages:

- every *fetch* goes to a different page of *S*
(this essentially returns us to the worst-case scenario for nested-loop join)
- the join only works in "one direction" (from *R* to *S*)
- requires additional data for each different join type
- requires tuples to be larger $\Rightarrow b_R$ is larger

General Join Conditions

83/87

Above examples all used simple equijoin e.g. $\text{Join}_{[i=j]}(R,S)$.

For theta-join e.g. $\text{Join}_{[i<j]}(R,S)$:

- index nested loop join: need B+ tree index on inner relation
- sort-merge join can be adapted, but is not very effective
- hash join is inapplicable
- other methods are essentially unchanged

... General Join Conditions

84/87

For multi-equality (pmr) join e.g. $\text{Join}_{[i=j \wedge k=l]}(R,S)$

- index nested loop join:
 - build index on all join fields of inner relation
 - e.g. if *S* is inner, build index on $(S.j, S.l)$
- sort-merge join:
 - sort both relations on combined join fields
 - e.g. sort *R* on $(R.i, R.k)$, sort *S* on $(S.j, S.l)$
- hash-join:
 - use multi-attribute hashing on combined join fields
- other methods are essentially unchanged

Join Summary

85/87

No single join algorithm is superior in some overall sense.

Which algorithm is best for a given query depends on:

- sizes of relations being joined, size of buffer pool
- any indexing on relations, whether relations are sorted
- which attributes and operations are used in the query
- number of tuples in *S* matching each tuple in *R*

Choosing the "best" join algorithm is critical because the cost difference between best and worst case can be very large.

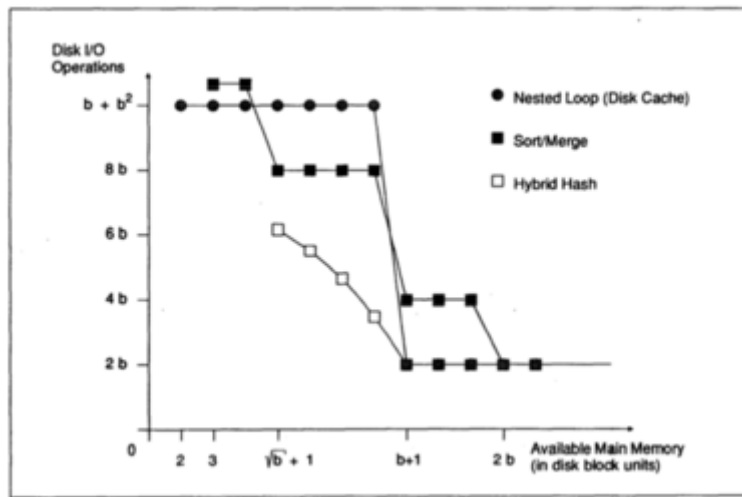
E.g. $\text{Join}_{[id=stude]}(Student, Enrolled)$: 3,000 ... 2,000,000

In some cases, it may be worth modifying access methods "on the fly" (e.g. add index) to enable an efficient join algorithm.

... Join Summary

86/87

Comparison of join costs (from Zeller/Gray VLDB90, assumes $b_R = b_S = b$)



Join in PostgreSQL

87/87

Join implementations are under: [src/backend/executor](#)

PostgreSQL supports three kinds of join:

- nested loop join ([nodeNestloop.c](#))
- sort-merge join ([nodeMergejoin.c](#))
- hash join ([nodeHashjoin.c](#)) (hybrid hash join)

Query optimiser chooses appropriate join, by considering

- physical characteristics of tables being joined
- estimated selectivity (likely number of result tuples)

Produced: 17 Jul 2019