

Exercises 01

Database Management Systems, PostgreSQL

Some of these questions require you to look beyond the Week 01 lecture material for answers. Some of the questions preempt material that we'll be looking at over the next few weeks. To answer some questions, you may need to look at the PostgreSQL documentation or at the texts for the course ... or, of course, you could simply reveal the answers, but where's the fun in that?

1. List some of the major issues that a relational database management system needs to concern itself with.

Answer:

- persistent storage of data and meta-data
 - executing SQL queries on stored data
 - maintenance of constraints on stored data
 - extensibility via views, triggers, procedures
 - query processing (optimisation and efficient execution)
 - transaction processing semantics
 - control of concurrent access by multiple users
 - recovery from failures (rollback, system failure)
2. Give an overview of the major stages in answering an SQL query in a relational database management system. For each step, describe its inputs and outputs and give a brief description of what it does.

Answer:

0. start with the text string of an SQL query
 - e.g. `select e.name,d.name from Employee e, Dept d where e.id=d.manager;`
 1. **parsing and translation**
 - converts an SQL query into a relational algebra expression
 - input: text string of an SQL query
 - output: expression tree for a relational algebra expression
 2. **optimisation**
 - converts RA expression into query plan
 - input: relational algebra expression tree
 - output: sequence of DBMS-specific relational operations
 3. **execution**
 - performs relational operations, via chained intermediate results
 - input: query plan (sequence of DBMS-specific relational operations)
 - output: set of result tuples, stored either in memory or on disk
 4. **output**
 - convert the result tuples into a format useful for the client
 - input: tuple data in memory buffers (and possibly on disk as well)
 - output: stream of formatted tuples (format defined by library e.g. PostgreSQL's `libpq`)
3. PostgreSQL is an "object-relational database management system". What are the differences between PostgreSQL and a "conventional" relational database management system (such as Oracle)?

Answer:

- every database tuple has an associated object identifier
- tables can be defined as specialisations of other tables (inheritance)

- can define new data types and operations on those types

4. A PostgreSQL installation includes a number of different "scopes": *databases* (or catalogs), *schemas* (or namespaces), and *tablespaces*. The scopes correspond to notions from the SQL standard. Explain the difference between these and give examples of each.

Answer:

- **database** (or **catalog**) ... a logical scope that collects together a number of schemas; an example is `template1`, a special database that is cloned whenever a user creates a new database; details of databases are held in the `pg_database` catalog table
- **schema** (or namespace) ... a logical scope used as a namespace; contains a collection of database objects (tables, views, functions, indexes, triggers, ...); an example is the `public` schema available as a default in all databases; details of schemas are held in the `pg_namespace` catalog table
- **tablespace** ... a physical scope identifying a region of the host filesystem where PostgreSQL data files are stored; an example is the `pg_default` tablespace, which corresponds to the `PG_DATA` directory where most PostgreSQL data files are typically stored; details of tablespaces are held in the `pg_tablespace` catalog table

5. For each of the following command-line arguments to the `psql` command, explain what it does, when it might be useful, and how you might achieve the same effect from within `psql`:

- `-l`
- `-f`
- `-a`
- `-E`

Answer:

a. `psql -l`

Generates a list of all databases in your cluster; would be useful if you couldn't remember the exact name of one of your databases.

You can achieve the same effect from within `psql` via the command `\list` or simply `\l`

b. `psql db -f file`

Connects to the database `db` and reads commands from the file called `file` to act on that database; useful for invoking scripts that build databases or that run specific queries on them; only displays the output from the commands in `file`.

You can achieve the same effect from within `psql` via the command `\i file`

c. `psql -a db -f file`

Causes all input to `psql` to be echoed to the standard output; useful for running a script on the database and being able to see error messages in the context of the command that caused the error.

You can achieve the same effect from within `psql` via the command `\set ECHO all`

d. `psql -E db`

Connect to the database `db` as usual; for all of the `psql` catalog commands (such as `\d`, `\df`, etc.), show the SQL query that's being executed to produce it; useful if you want to learn how to use the catalog tables.

You can achieve the same effect from within `psql` via the command `\set ECHO_HIDDEN on`

6. PostgreSQL has two main mechanisms for adding data into a database: the SQL standard `INSERT` statement and the PostgreSQL-specific `COPY` statement. Describe the differences in how these two statements operate. Use the following examples, which insert the same set of tuples, to motivate your explanation:

```
insert into Enrolment(course,student,mark,grade)
    values ('COMP9315', 3312345, 75, 'DN');
insert into Enrolment(course,student,mark,grade)
    values ('COMP9322', 3312345, 80, 'DN');
insert into Enrolment(course,student,mark,grade)
    values ('COMP9315', 3354321, 55, 'PS');

copy Enrolment(course,student,mark,grade) from stdin;
COMP9315      3312345 75      DN
COMP9322      3312345 80      DN
COMP9315      3354321 55      PS
\.
```

Answer:

Each `insert` statement is a transaction in its own right. It attempts to add a single tuple to the database, checking all of the relevant constraints. If any of the constraints fails, that particular insertion operation is aborted and the tuple is not inserted. However, any or all of the other `insert` statements may still succeed.

A `copy` statement attempts to insert all of the tuples into the database, checking constraints as it goes. If any constraint fails, the `copy` operation is halted, and none of the tuples are added to the table†.

For the above example, the `insert` statements may result in either zero or 1 or 2 or 3 tuples being inserted, depending on whether how many values are valid. For the `copy` statement, either zero or 3 tuples will be added to the table, depending on whether any tuple is invalid or not.

† A fine detail: under the `copy` statement, tuples are "temporarily" added to the table as the statement progresses. In the event of an error, the tuples are all marked as invalid and are not visible to any query (i.e. they are effectively *not* added to the table). However, they still occupy space in the table. If a very large `copy` loads e.g. 9999 or 10000 tuples and the last tuple is incorrect, space has still been allocated for the most of the tuples. The `vacuum` function can be used to clean out the invalid tuples.

7. In `psql`, the `\timing` command turns on a timer that indicates how long each SQL command takes to execute. Consider the following trace of a session asking the several different queries multiple times:

```
unsw=# \timing
Timing is on.
unsw=# select max(id) from students;
      max
-----
9904944
Time: 112.173 ms
unsw=# select max(id) from students;
      max
-----
9904944
Time: 0.533 ms
```

```

unsw=# select max(id) from students;
      max
-----
9904944
Time: 0.484 ms
unsw=# select count(*) from courses;
      count
-----
80319
Time: 132.416 ms
unsw=# select count(*) from courses;
      count
-----
80319
Time: 30.438 ms
unsw=# select count(*) from courses;
      count
-----
80319
Time: 34.034 ms
unsw=# select max(id) from students;
      max
-----
9904944
Time: 0.765 ms
unsw=# select count(*) from enrolments;
      count
-----
2816649
Time: 2006.707 ms
unsw=# select count(*) from enrolments;
      count
-----
2816649
Time: 1099.993 ms
unsw=# select count(*) from enrolments;
      count
-----
2816649
Time: 1109.552 ms

```

Based on the above, suggest answers to the following:

- Why is there such variation in timing between different executions of the same command?
- What timing value should we ascribe to each of the above commands?
- How could we generate reliable timing values?
- What is the accuracy of timing results that we can extract like this?

Answer:

a. Variation:

There's a clear pattern in the variations: the first time a query is executed it takes *significantly* longer than the second time its executed (e.g. the first query drops from over 100ms to less than 1ms). This is due to caching effects. PostgreSQL has a large in-memory buffer-pool. The first time a query is executed, the relevant pages will need to be read into memory buffers from disk. The second and subsequent times, the pages are already in the memory buffers.

b. Times:

Given the significantly different contexts, it's not really plausible to assign a specific time to a query. Assigning a range of values, from "cold" execution (when none of the data for the

query is buffered) to "hot" execution (when as much as possible of the needed data is buffered), might be more reasonable. Even then, you would need to measure the hot and cold execution several times and take an average.

How to achieve "cold" execution multiple times? It's difficult. Even if you stop the PostgreSQL server, then restart it, effectively flushing the buffer pool, there is still some residual buffering in the Unix file buffers. You would need to read lots of other files to flush the Unix buffers.

c. Reliability:

This is partially answered in the previous question. If you can ensure that the context (hot or cold) is the same at the start of each timing, the results will be plausibly close. Obviously, you should run each test on the same lightly-loaded machine (to minimise differences caused by Unix buffering). You should also ensure that you are the only user of the database server. If multiple users are competing for the buffer pool, the times could variably substantially and randomly up or down between subsequent runs, depending on how much of your buffered data had been swapped out to service queries from other users.

d. Accuracy:

For comparable executions of the query (either buffers empty or buffers fully-loaded), it looks like it's no more accurate than $\pm 10\text{ms}$. It might even be better to forget about precise time measures, and simply fit queries into "ball-park" categories, e.g.

- Very fast ... $0 \leq t < 100\text{ms}$
- Fast ... $100 \leq t < 500\text{ms}$
- Acceptable ... $500 \leq t < 2000\text{ms}$
- Slow ... $2000 \leq t < 10000\text{ms}$
- Too Slow ... $t > 10000\text{ms}$

Note that the above queries were run on a PostgreSQL 8.3.5 server. More recent servers seem to be somewhat more consistent in the value returned for "hot" executions, although there is may still be a substantial difference between the first "cold" execution of a query and subsequent "hot" executions of the same query.